Global PC Component Price Intelligence Platform

Product Plan and Roadmap

Working concept: A global market-intelligence and retail-price platform explaining how semiconductor-market events affect the price consumers pay for PC components in different regions.

Working proposition:

> One global component market. Very different local prices.

Core promise:

> We track global semiconductor-market pressure, measure how and when it reaches regional retail prices, explain why prices are moving, and publicly score the accuracy of our forecasts.

---

1. Executive summary

The platform will combine four forms of information:

1. **Global market intelligence**

2. DRAM and NAND supply

3. AI and HBM demand
4. Manufacturer production decisions
5. Data-centre investment
6. Product and process transitions

7. Trade, tariff and currency events

8. **Regional retail-price tracking**

9. Current component prices

10. Historical price movement
11. Stock availability
12. New and used-market comparisons

13. Regional and retailer differences

14. **Editorial analysis**

15. Why prices moved

16. How strongly an event is likely to affect consumers
17. Expected delay before upstream changes reach retail

18. Evidence supporting and opposing each conclusion

19. **Consumer guidance**

20. Buy now, wait or neutral

21. Regional price outlooks
22. Value alternatives
23. Complete-build cost impact
24. Forecast confidence and past accuracy

The site will be global from the start, but retail-price coverage will launch in selected markets.

The initial recommended markets are:

- United States
- Eurozone, initially represented primarily by Germany
- United Kingdom
- Japan

The initial recommended component categories are:

- DDR4 desktop memory
- DDR5 desktop memory
- Consumer NVMe SSDs

The platform should be designed to add further countries and product categories without rebuilding the underlying system.

---

2. Product vision

2.1 The problem

PC buyers currently receive fragmented information.

Price-comparison platforms can show the current price of a product. Historical-price tools can show whether a listing was cheaper previously. Technology publications report manufacturing announcements and shortages. Market-research organisations discuss contract pricing and supply.

Very few services connect all of these layers.

A buyer may see:

- DRAM contract prices increasing
- A technology article claiming shortages are worsening
- A retailer discounting a particular memory kit
- Someone online saying prices are already falling

All four observations may be true simultaneously.

The missing explanation is that:

- Upstream prices and retail prices move at different speeds.
- Retailers hold different amounts of older inventory.
- Exchange rates affect regional prices.
- Particular products may be discounted while the overall market rises.
- Contract, spot, wholesale and retail prices are different measures.

The platform will make those distinctions understandable.

2.2 The intended position

The platform should not position itself as another component catalogue or generic price-comparison site.

It should become:

> The consumer-facing intelligence layer between semiconductor manufacturing and the retail checkout.

2.3 Long-term ambition

The long-term business can develop into three connected products:

Consumer publication

Free price tracking, market explanations, news analysis, recommendations and buying guidance.

Premium consumer product

Alerts, watchlists, detailed histories, advanced forecasts, regional comparisons and complete-build tracking.

Professional market-intelligence product

Data feeds, dashboards, reports and APIs for:

- PC retailers
- System integrators

- Technology journalists
- Reviewers and content creators
- Procurement teams
- Repair businesses
- Analysts and investors
- Component manufacturers

---

3. Unique selling proposition

3.1 Primary USP

> The platform measures how global semiconductor-market changes pass through into regional consumer prices.

The product will not simply report that DRAM became more expensive. It will attempt to show:

- What caused the upstream pressure
- Which component classes are affected
- Which regions have already reacted
- Which regions appear insulated by existing inventory
- How exchange rates affect the result
- How long the retail response may take
- Whether the consumer should buy or wait

3.2 Retail Price Lag

The signature analytical feature should be a **Retail Price Lag** indicator.

This measures the apparent delay between an upstream market event and a material change in regional retail prices.

Example:

Layer	Status
Global DRAM pressure	Strongly upward
US retail response	Rising rapidly
German retail response	Rising moderately
UK retail response	Delayed
Japanese retail response	Amplified by currency

Initially, this can be a transparent editorial classification.

Later, it can become a calculated metric using:

- Upstream index movement
- Regional retail-index movement
- Exchange-rate changes
- Retail inventory indicators
- Time-series correlation
- Known product and distribution delays

3.3 Forecast accountability

Every published forecast should remain permanently accessible.

Each forecast will record:

- Publication date
- Region
- Product category
- Time horizon
- Expected direction
- Expected percentage range
- Confidence level
- Supporting evidence
- Contradictory evidence
- Conditions that would invalidate the forecast
- Review date
- Actual result

Example:

Forecast published: 5 August 2026
Market: United States
Category: 32GB DDR5 desktop kits
Horizon: 90 days
Forecast: Median retail price likely to rise by 5–15%
Confidence: 65%
Outcome: Recorded automatically after 90 days

The public scorecard should include:

- Directional accuracy
- Average forecast error
- Accuracy by component
- Accuracy by region
- Accuracy by forecast horizon
- Confidence calibration
- Correct, partially correct and incorrect forecasts

This will distinguish the platform from publications that make predictions without revisiting them.

---

4. Target audiences

4.1 Primary consumer audiences

PC builders

People planning a complete new gaming, workstation or enthusiast system.

Their main question:

> Should I buy these components now, or is waiting likely to save me money?

Upgraders

People replacing RAM, storage, a GPU or another individual part.

Their main question:

> Is the current price reasonable relative to recent history?

Deal-focused buyers

People actively looking for components below their typical market value.

Their main question:

> Is this actually a deal, or is the displayed discount misleading?

Technology enthusiasts

People interested in component markets, manufacturing, shortages and industry developments.

Their main question:

> Why is this happening, and what is likely to happen next?

4.2 Professional audiences

- Independent PC retailers
- System builders
- IT procurement teams
- Review publications

- YouTube creators
- Technology journalists
- Repair companies
- Market analysts
- Manufacturers and distributors

These users will eventually need:

- Downloadable data
- Regional comparisons
- Stock-pressure indicators
- Charts licensed for publication
- API access
- Category forecasts
- Custom reporting

---

5. Product architecture

The platform should be divided into three logical layers.

5.1 Global intelligence layer

This layer is shared across all regions.

It will include:

- Global DRAM outlook
- Global NAND outlook
- Memory-manufacturer developments
- HBM and AI demand
- Manufacturing capacity
- Production cuts and expansions
- Process and technology transitions
- Data-centre capital expenditure
- Worldwide consumer-device demand
- Trade restrictions and tariffs
- Global forecast
- Industry-event timeline

The default reference currency for upstream information should be USD.

5.2 Regional retail layer

This layer shows the consumer-facing outcome.

Each region will have:

- Local currency
- Regional retailers
- Regional affiliate links
- Relevant tax treatment
- Regional retail indices
- Stock availability
- Local price outlook
- Regional event impact
- Used-market information where available
- Regional buying guidance

5.3 Decision layer

This layer converts data into practical advice.

It will show:

- Buy now
- Consider buying
- Neutral
- Consider waiting
- Wait where practical

Every conclusion must show:

- Time horizon
- Confidence
- Main reasons
- Risks to the conclusion
- Date last reviewed
- Forecast history

The recommendation should never be presented as certainty.

---

6. Geographic strategy

6.1 Initial launch markets

United States

Reasons:

- Large PC-building audience

- USD reference market
- Strong affiliate and advertising potential
- Broad retailer coverage
- Large volume of technology search traffic

Germany and Eurozone overview

Reasons:

- Major European PC-hardware market
- EUR reference market
- Strong specialist-retailer ecosystem
- Useful foundation for wider European expansion

Germany should initially provide much of the data behind a clearly labelled Eurozone retail indicator. It should not be presented as perfectly representative of every euro-area country.

United Kingdom

Reasons:

- Distinct currency
- Strong enthusiast market
- Good specialist retailers
- Useful comparison with the Eurozone
- Existing team knowledge

Japan

Reasons:

- Major technology market
- JPY currency effects
- Important Asian comparison
- Strong contrast with US and European retail behaviour

6.2 Expansion markets

Second wave

- Canada
- France
- Netherlands
- Italy
- Spain
- Australia
- India
- Singapore

Third wave

- Poland
- Nordic countries
- South Korea
- Taiwan
- Hong Kong
- China
- Mexico
- Brazil and selected Latin American markets

6.3 Geographic expansion criteria

A new country should only be added when the team can provide:

- At least three reliable retail sources
- Stable product matching
- Correct tax and currency treatment
- Local affiliate or commercial potential
- Sufficient recurring user demand
- Someone capable of validating regional anomalies
- A documented regional methodology

A region should not be marked as fully supported merely because one retailer can be queried.

---

7. Currency and tax methodology

7.1 Global upstream data

Use USD for:

- Contract-price information
- Spot-price information
- Manufacturer financial data
- Market-size information
- Data-centre spending
- Global forecasts
- Cross-regional normalisation

7.2 Regional retail data

Use local currencies:

- United States: USD
- Canada: CAD

- United Kingdom: GBP
- Eurozone: EUR
- Japan: JPY
- Australia: AUD
- India: INR

7.3 Regional comparison modes

Every regional chart should eventually support:

Local price

The price actually relevant to a buyer in that country.

USD-equivalent price

Useful for cross-regional analysis.

Indexed movement

Rebase each regional basket to 100 at the beginning of the selected period.

This is the preferred method for comparing how quickly regional prices are rising or falling.

7.4 Tax handling

Prices must be clearly labelled as:

- Tax included
- Tax excluded
- Estimated tax
- Tax varies by location

US prices should not be compared directly with European VAT-inclusive prices without a clear adjustment or warning.

The site must retain both:

- Raw retailer price
- Normalised comparison price

---

8. Initial product coverage

8.1 Launch categories

DDR4 desktop memory

Recommended baskets:

- 16GB DDR4-3200
- 32GB DDR4-3200
- 64GB DDR4-3200

DDR5 desktop memory

Recommended baskets:

- 32GB DDR5-5600
- 32GB DDR5-6000
- 64GB DDR5-6000
- Higher-speed enthusiast kit basket where data allows

NVMe SSDs

Recommended baskets:

- 1TB PCIe 4.0
- 2TB PCIe 4.0
- 4TB PCIe 4.0
- Premium versus value-tier comparison

8.2 Excluded from the first release

- Laptop memory
- Server DIMMs
- ECC memory
- Enterprise SSDs
- SATA SSDs unless a clear audience case emerges
- GPUs
- CPUs
- Motherboards
- Complete PCs
- Highly specialised overclocking products

These can be added after the data and methodology are stable.

8.3 Product-selection rules

Products included in an index should meet documented standards.

Suggested requirements:

- Sold by a reputable retailer
- In stock or orderable
- Clearly identifiable manufacturer part number
- Correct capacity and specification
- Not a marketplace listing from an unknown third party
- Not an obvious clearance anomaly unless specifically labelled
- Not bundled with unrelated products
- Not refurbished unless part of a refurbished index
- Available in sufficient quantity or across sufficient retailers

---

9. Price-index methodology

9.1 Why category indices are needed

Individual-product prices can be distorted by:

- Temporary promotions
- Product discontinuation
- Low stock
- Marketplace sellers
- Replacement models
- Incorrect listings
- Bundles
- Retailer pricing mistakes

The primary market indicator should therefore be a category index rather than one product.

9.2 Proposed calculation

For each regional basket:

1. Collect qualifying product offers.
2. Match offers to canonical product records.
3. Remove duplicates and invalid listings.
4. Exclude extreme anomalies according to published rules.
5. Calculate the median qualifying retail price.
6. Retain minimum, maximum and product-count values.
7. Save one official daily index value.

8. Record changes in product composition.
9. Flag large movements for human review.

The median is recommended initially because it is less sensitive to extreme prices than the mean.

9.3 Index quality score

Each index should show a data-quality indicator based on:

- Number of products
- Number of retailers
- Percentage of products successfully matched
- Percentage currently in stock
- Age of latest update
- Presence of abnormal values

Possible ratings:

- High confidence
- Moderate confidence
- Limited confidence
- Insufficient data

9.4 Product substitutions

When a product disappears:

- Do not silently replace it.
- Record the retirement date.
- Add its successor as a new canonical product.
- Document whether the basket composition changed.
- Avoid allowing one brand to dominate an index.

9.5 Historical revisions

Raw price snapshots should never be overwritten.

Corrections should create:

- Original value
- Corrected value
- Reason for correction
- Timestamp
- Editor or automated process responsible

This audit trail will be important for trust and future commercial licensing.

---

10. News and event methodology

10.1 Event types

Every event should be tagged as one or more of:

- Supply increase
- Supply reduction
- HBM or AI demand
- Server demand
- Consumer demand
- Inventory correction
- Manufacturing disruption
- New fabrication capacity
- Process transition
- Product launch
- Currency movement
- Tariff
- Export restriction
- Regulatory change
- Retail promotion
- Distribution issue
- Natural disaster
- Earnings guidance
- Capital expenditure

10.2 Geographic scope

Each event should be marked:

- Global
- North America
- Europe
- Asia-Pacific
- Country-specific

10.3 Source hierarchy

Use a visible source classification.

Tier 1: Primary sources

- Manufacturer statements
- Investor presentations
- Earnings reports
- Regulatory filings

- Government announcements
- Industry standards organisations

Tier 2: Specialist market research

- Semiconductor research organisations
- Supply-chain analysis firms
- Industry associations

Tier 3: Major financial reporting

- Reputable financial and business publications
- News agencies
- Established newspapers

Tier 4: Specialist technology reporting

- Established PC-hardware and semiconductor publications

Tier 5: Retail and community signals

- Retailer stock information
- Forum discussion
- Social-media reports
- User submissions

Tier 5 evidence should be treated as a signal requiring confirmation, not as proof.

10.4 Event-analysis template

Every major event should answer:

1. What happened?
2. Which products or markets could be affected?
3. Is the expected pressure upward, downward or uncertain?
4. How strong could the effect be?
5. When might consumers see the effect?
6. Which regions are most exposed?
7. What evidence supports this?
8. What evidence argues against it?
9. What would change the conclusion?
10. When will the assessment be reviewed?

10.5 Editorial standards

The site must:

- Link to original sources

- Use original wording for analysis
- Separate facts from interpretation
- Clearly identify uncertainty
- Correct errors publicly
- Avoid automatically generated causal claims
- Avoid republishing copyrighted articles
- Mark sponsored content prominently
- Prevent advertisers from influencing indices or forecasts

---

11. Forecasting framework

11.1 Initial forecast model

The first version should be editorial and rules-based rather than pretending to have a sophisticated predictive model.

Inputs can include:

- Recent retail-index movement
- Upstream contract-price direction
- Spot-price direction
- Manufacturer guidance
- Capacity announcements
- Inventory indications
- Currency changes
- Seasonal demand
- Product-launch cycle
- Retail stock levels
- Regional promotion patterns

11.2 Forecast output

Each forecast should include:

- Region
- Category
- Current index value
- One-month direction
- Three-month direction
- Six-month direction where justified
- Expected movement range
- Confidence
- Primary drivers
- Counterarguments
- Invalidation conditions

- Review date

11.3 Direction labels

Recommended labels:

- Strongly rising
- Rising
- Broadly stable
- Falling
- Strongly falling
- Highly uncertain

11.4 Consumer recommendation labels

Recommended labels:

- Buy now
- Consider buying
- Neutral
- Consider waiting
- Wait where practical

The recommendation must consider more than expected price direction.

For example, a rising market may still receive “neutral” where:

- Current prices are already unusually high
- A replacement product is imminent
- Availability remains good
- The expected increase is small

11.5 Forecast scorecard

The initial scorecard should measure:

- Correct direction
- Actual percentage change
- Forecast midpoint error
- Whether the actual outcome fell within the forecast range
- Forecast confidence
- Data quality at the time of prediction

---

12. Core website experience

12.1 Homepage

The homepage should answer five questions immediately:

1. Are global memory prices rising or falling?
2. Why?
3. Which regions are being hit hardest?
4. What should a consumer do?
5. How confident is the platform?

Suggested homepage order:

Global market status

- DRAM pressure
- NAND pressure
- Overall direction
- Three-month outlook
- Last updated

Regional transmission

A comparison of:

- United States
- Eurozone
- United Kingdom
- Japan

Main price charts

- 32GB DDR5
- 64GB DDR5
- 2TB NVMe SSD

What moved the market

The three to five most important recent events.

Buy or wait

Current recommendations by category and region.

Forecast scorecard

Headline accuracy metrics.

Current opportunities

Products materially below their regional average.

12.2 Global market page

Contains:

- Global market outlook
- Upstream indicators
- Manufacturer timeline
- Major demand drivers
- Global events
- Regional comparison
- Methodology

12.3 Regional page

Contains:

- Local-currency indices
- USD-equivalent view
- Indexed comparison
- Regional retailers
- Currency pressure
- Stock conditions
- Regional forecast
- Regional affiliate links
- Relevant local events

12.4 Category page

Example: DDR5 memory.

Contains:

- Global category outlook
- Regional comparison
- Capacity and speed filters
- Historical index
- Major events shown on the graph
- Buy/wait recommendation

- Current products below average
- Forecast archive

12.5 Product page

Contains:

- Current qualifying prices
- Regional retailer links
- Price history
- Position versus category median
- Lowest recorded price
- Product specifications
- Availability
- Closely matched alternatives
- Affiliate disclosure

12.6 Event page

Contains:

- Event summary
- Primary source
- Supporting sources
- Expected global impact
- Regional impact
- Components affected
- Timing estimate
- Confidence
- Related forecasts
- Subsequent updates

12.7 Methodology page

This is essential.

It should explain:

- Product selection
- Retailer selection
- Tax handling
- Currency conversion
- Index calculation
- Outlier handling
- Product retirement
- Event classification
- Forecast scoring

- Affiliate relationships
- Corrections policy

---

13. Technical architecture

13.1 Main system components

Data ingestion

Collects:

- Retailer feeds
- Affiliate feeds
- Permitted APIs
- Currency data
- Used-market listings
- Manually entered events
- Manufacturer and industry sources

Product normalisation

Converts retailer listings into canonical products.

Price history

Stores immutable raw observations and derived daily values.

Index engine

Calculates regional category indices.

Editorial system

Allows the team to:

- Add events
- Attach evidence
- Assign impact
- Publish analysis
- Create forecasts
- Review forecast outcomes
- Correct data

Public API layer

Supplies the website with:

- Prices
- Indices
- Events
- Forecasts
- Regional settings
- Product data

Analytics and monitoring

Tracks:

- Failed imports
- Missing prices
- Abnormal movements
- Product-match confidence
- Affiliate clicks
- Search performance
- User retention

13.2 Recommended data entities

- Region
- Country
- Currency
- Retailer
- Affiliate programme
- Manufacturer
- Canonical product
- Retailer listing
- Product specification
- Raw price observation
- Daily product price
- Category basket
- Daily index
- Exchange rate
- Event
- Source
- Event impact
- Forecast
- Forecast result
- Editorial correction
- Affiliate click
- User watchlist

- Alert

13.3 Product matching

Product matching should prioritise:

1. Manufacturer part number
2. Barcode or GTIN
3. Manufacturer
4. Capacity
5. Memory type
6. Speed
7. Latency
8. Kit configuration
9. Colour only where technically relevant
10. Retailer title as a lower-confidence signal

Every automated match should receive a confidence score.

Low-confidence matches should enter a human-review queue.

13.4 Data-quality monitoring

Automatic alerts should trigger when:

- A price changes by more than a defined threshold
- Product count drops sharply
- A retailer feed fails
- Currency data is stale
- A product changes specification
- A regional index moves unexpectedly
- A marketplace seller becomes the cheapest offer
- Stock disappears from multiple retailers
- Product matching confidence falls

---

14. Team structure and ownership

The following roles can be combined in a small team.

Product owner

Owns:

- Product strategy

- Scope
- Prioritisation
- Commercial model
- Release decisions
- Stakeholder alignment

Technical lead

Owns:

- Architecture
- Engineering standards
- Security
- Deployment
- Reliability
- Technical roadmap

Data engineer

Owns:

- Retailer ingestion
- APIs
- Product matching
- Price history
- Index calculations
- Data-quality monitoring

Front-end developer or product designer

Owns:

- User experience
- Charts
- Regional navigation
- Mobile experience
- Accessibility
- Design system

Market editor or analyst

Owns:

- Source selection
- Event analysis
- Forecasts

- Corrections
- Methodology
- Editorial standards

Growth and commercial lead

Owns:

- SEO
- Newsletter
- Affiliate programmes
- Sponsorship
- Retail partnerships
- B2B customer discovery

Legal and privacy adviser

Part-time or external responsibility for:

- Privacy
- Cookies
- Affiliate disclosures
- Data licensing
- Terms of service
- Paid subscriptions
- Sponsored content

---

15. Delivery roadmap

Phase 0: Alignment and validation

Weeks 1–2

Objectives

- Confirm product position
- Agree launch scope
- Test data availability
- Assign ownership
- Define success criteria

Work

Product

- Confirm working brand and proposition.
- Define initial target users.
- Agree the first four regions.
- Agree initial product baskets.
- Document what is explicitly out of scope.
- Create a competitor matrix.
- Conduct 10–15 user interviews.

Suggested interview groups:

- PC builders
- Enthusiasts
- Retail staff
- System integrators
- Technology writers
- YouTube creators

Data

- Identify three or more viable price sources per launch region.
- Test product-feed formats.
- Test eBay or other used-market sources.
- Confirm manufacturer-part-number coverage.
- Sample 50 listings and estimate matching difficulty.
- Establish tax and currency rules.

Editorial

- Build the initial source list.
- Define event tags.
- Draft the forecast template.
- Produce two example market analyses.
- Establish source and corrections policies.

Commercial

- Identify affiliate programmes by region.
- Check approval requirements.
- Build a list of potential launch partners.
- Interview at least three possible B2B users.

Deliverables

- Product requirements document
- Competitor matrix

- Data-source register
- Initial methodology
- User-interview summary
- Risk register
- Launch scope
- Team ownership matrix

Exit criteria

Proceed only when:

- At least three regions have workable retail data.
- At least 70% of the sample product listings can be matched reliably.
- Users understand and value the global-to-regional proposition.
- The team can explain the product in one sentence.
- No critical data-licensing issue remains unresolved.

Phase 1: Data and editorial foundations

Weeks 3–6

Objectives

- Establish the canonical product database
- Begin collecting daily price history
- Build the editorial workflow
- Prove one regional index end to end

Engineering work

- Create region, retailer and product schemas.
- Create raw price-observation storage.
- Build the first retailer-feed import.
- Build canonical product matching.
- Create manual product-review tools.
- Create exchange-rate ingestion.
- Create daily index calculation.
- Create anomaly detection.
- Create data-quality score.
- Create basic internal dashboard.

Editorial work

- Build the event-entry interface.
- Add source classification.
- Add event-impact fields.

- Create forecast records.
- Create review dates.
- Create corrections workflow.
- Draft methodology pages.

Initial proof market

Use the United States or whichever region has the cleanest available source data.

Prove:

- One complete 32GB DDR5 index
- Daily price collection
- Product matching
- Index generation
- Event publication
- One forecast
- One product page

Deliverables

- Working product catalogue
- Working price ingestion
- First daily category index
- Internal data-quality dashboard
- Editorial event workflow
- Forecast workflow
- Methodology draft

Exit criteria

- Seven consecutive successful daily imports
- No unresolved critical data corruption
- Index can be reproduced from raw records
- More than 85% of included offers are correctly matched
- Editors can publish without developer assistance
- Every public value has a traceable source

Phase 2: MVP build

Weeks 7–12

Objectives

- Complete the public MVP
- Add the four launch regions

- Add core product categories
- Prepare private beta

Public features

- Global homepage
- Regional selector
- United States page
- Eurozone or Germany page
- United Kingdom page
- Japan page
- DDR4 category page
- DDR5 category page
- NVMe SSD category page
- Product pages
- Event pages
- Price charts
- Regional comparison
- Buy/wait recommendations
- Forecast archive
- Methodology
- Affiliate disclosure
- Privacy and cookie controls
- Newsletter signup

Data targets

- 30–60 canonical products
- 10–15 meaningful category baskets
- At least three reliable retail sources in each fully supported region
- Daily price updates
- Daily currency updates
- Initial used-market indicator in at least one region

Editorial targets

Before private beta, publish:

- One global memory-market overview
- Four regional outlooks
- Three category explainers
- Ten major historical or current event entries
- Three active forecasts
- A forecast-scoring explanation
- A clear corrections policy

Quality assurance

Test:

- Mobile usability
- Regional currency switching
- Tax labels
- Broken affiliate links
- Out-of-stock behaviour
- Duplicate products
- Incorrect specifications
- Chart readability
- Accessibility
- Search indexing
- Page performance

Exit criteria

- Price data remains stable for 30 consecutive days.
- Every supported index shows its data-quality status.
- No region has fewer than the stated minimum data sources.
- Every forecast displays supporting and contradictory evidence.
- Every affiliate link is disclosed.
- All major pages work on mobile.
- The team can complete the weekly editorial workflow within the planned labour budget.

Phase 3: Private beta

Months 4–5

Objectives

- Test whether users return
- Validate recommendations
- Find data-quality weaknesses
- Improve the editorial voice

Beta group

Recruit 100–300 users from:

- PC-building communities
- Friends and professional contacts
- Technology forums
- Retail and system-building contacts

- YouTube and social audiences

Beta activities

- Weekly feedback survey
- User interviews
- Task testing
- Monitor regional selection
- Measure chart engagement
- Measure buy/wait usage
- Review affiliate clicks
- Log every data complaint
- Test newsletter format
- Publish weekly forecast updates

Primary questions

- Do users understand the difference between global and local prices?
- Do users trust the recommendation?
- Is confidence presented clearly?
- Which charts are genuinely useful?
- Are regional comparisons understandable?
- Do users return weekly?
- Do users click through to retailers?
- Which categories generate the most interest?

Beta targets

- 25% or more of users return within 30 days.
- 15% or more subscribe to the newsletter.
- Fewer than 2% of viewed product pages generate a data-quality complaint.
- At least 50 users use the regional comparison more than once.
- At least 20 users report changing or confirming a buying decision.

Deliverables

- Beta findings report
- Revised homepage
- Revised methodology
- Data-source quality ranking
- Prioritised launch backlog
- First public forecast outcomes

Phase 4: Public launch

Month 6

Launch package

Publish:

- Global market state
- Regional price comparison
- “How global component prices reach local retailers”
- Current RAM and SSD outlook
- First forecast scorecard
- Original launch report
- Full methodology
- Newsletter issue
- Media and creator briefing pack

Launch outreach

Target:

- PC-building publications
- Technology journalists
- YouTube reviewers
- Hardware forums
- Reddit communities
- Discord communities
- Retail and system-building partners
- Semiconductor and data-centre commentators

Launch content angle

The strongest launch asset should be an original report such as:

> How quickly global memory-price shocks reach consumers in the US, Europe, Britain and Japan.

The report should contain original charts that publications and creators can cite.

Launch success measures

- 10,000 monthly visits within the first three months
- 1,000 newsletter subscribers
- At least five quality inbound links
- At least three press or creator mentions

- First affiliate conversions
- At least one sponsor discussion
- Forecast scorecard populated with real outcomes

Phase 5: Growth and reinvestment

Months 7–12

Product expansion

Add based on demand:

- Canada
- France
- Netherlands
- Australia
- India
- Singapore

Add selected product categories:

- GPUs
- CPUs
- Complete-build cost index
- Prebuilt gaming PCs

Do not add all categories simultaneously.

User features

- Price alerts
- Product watchlists
- Regional default
- Saved builds
- Complete-build cost tracking
- Weekly personalised digest
- Used-versus-new comparison
- Historical forecast filters

Editorial growth

- Weekly global outlook
- Regional monthly reports
- Earnings-season coverage
- Manufacturer profiles
- Retail-lag investigations

- Quarterly forecast review
- Data-led original reports

Commercial growth

- Improve affiliate attribution.
- Introduce restrained display advertising.
- Sell newsletter sponsorship.
- Test supporter membership.
- Interview professional users.
- Create sample B2B reports.
- Pilot chart licensing.
- Approach retailers for direct feeds.

Reinvestment gate

Begin paid tools and freelance support only after recurring revenue or clear traffic justifies them.

Suggested threshold:

- At least £500–£1,000 monthly revenue for three consecutive months

Priority order for reinvestment:

1. Better retail data
2. Product-data cleaning
3. Original editorial work
4. Newsletter tooling
5. SEO research
6. Limited AI assistance
7. Commercial research data

Phase 6: Commercial intelligence platform

Months 12–24

Objectives

- Develop professional subscriptions
- Introduce licensed datasets
- Build an API
- Add direct commercial partnerships
- Expand geographic and product coverage

Professional product

Possible features:

- Regional component dashboards
- Wholesale-to-retail spread
- Retail Price Lag metrics
- Stock-pressure indicators
- Product retirement and replacement alerts
- Downloadable CSV
- Forecast reports
- Data API
- Commercial chart licences
- Scheduled reports
- Retailer benchmarking

Commercial validation

Before major investment, secure:

- At least three design-partner businesses
- At least one paid data pilot
- Written feedback on required metrics
- A clear pricing test
- Evidence that the data affects business decisions

Suggested professional tiers

Builder Pro

For small system builders and repair companies.

- Category forecasts
- Stock pressure
- Regional trends
- Downloadable reports

Retail Intelligence

For retailers and distributors.

- Competitor pricing
- Product-level history
- Regional comparison
- Replacement detection
- Alerts

Data and API

For media, analysts and larger businesses.

- API access
- Bulk downloads
- Chart licensing
- Custom feeds
- Historical datasets

---

16. First 90-day work plan

Weeks 1–2

Product

- Finalise proposition
- Select working brand
- Define launch audience
- Confirm supported regions
- Confirm initial indices
- Interview users
- Create competitor map

Data

- Evaluate feeds and APIs
- Create sample product catalogue
- Test 50–100 retailer listings
- Confirm currency and tax handling
- Document source licences

Editorial

- Build source register
- Create event taxonomy
- Draft sample global analysis
- Draft sample regional analysis
- Define forecast format

Weeks 3–4

Engineering

- Set up database

- Create product and retailer models
- Create raw-price table
- Build first feed import
- Begin daily snapshots
- Create matching-confidence system

Editorial

- Create editorial templates
- Draft methodology
- Create first ten events
- Create first experimental forecast

Design

- Wireframe homepage
- Wireframe regional page
- Wireframe category page
- Wireframe event page

Weeks 5–6

Engineering

- Build first index
- Add anomaly detection
- Add exchange-rate data
- Add internal review queue
- Build public read API

Product

- Review index accuracy
- Finalise supported terminology
- Approve chart presentation
- Define confidence labels

Commercial

- Apply to affiliate programmes
- Build partner list
- Prepare disclosure language
- Start newsletter account

Weeks 7–8

Engineering

- Add further regions
- Add chart endpoints
- Build forecast records
- Build event-to-region relationships
- Add regional currency display

Front end

- Build homepage
- Build region selector
- Build first regional page
- Build category chart
- Build methodology page

Editorial

- Publish internal global outlook
- Publish internal regional outlooks
- Review source quality

Weeks 9–10

Engineering

- Add remaining MVP regions
- Add product pages
- Add affiliate tracking
- Add data-quality badge
- Add alerting for failed imports

Editorial

- Complete launch articles
- Complete event timeline
- Add contradictory evidence to forecasts
- Review all terminology

QA

- Test tax labels
- Test currencies
- Test product matching
- Test out-of-stock rules
- Test chart scales

Weeks 11–12

Private beta preparation

- Recruit beta users
- Create feedback survey
- Add analytics
- Test newsletter
- Finalise privacy and affiliate disclosures
- Fix high-priority defects
- Freeze MVP scope
- Begin private beta

---

17. Product backlog priorities

Must have for MVP

- Global market overview
- Four regional views
- Local currency
- USD comparison
- Indexed regional comparison
- Daily price history
- Category indices
- Event timeline
- Event impact analysis
- Buy/wait outlook
- Forecast archive
- Methodology
- Data-quality labels
- Affiliate disclosure
- Newsletter signup
- Mobile support

Should have shortly after launch

- Price alerts
- Product watchlists
- User region preference
- Used-market indicator
- Complete-build cost index
- Better forecast scorecard
- Automated forecast reviews
- Retailer stock trend
- Downloadable public charts

Could have later

- User accounts
- Saved builds
- Paid membership
- Professional dashboard
- API
- Custom reports
- Retailer benchmarking
- Mobile application
- Machine-learning forecast model

Explicitly avoid initially

- Covering every component category
- Publishing AI-generated articles automatically
- Paying for broad news APIs
- Building a social network
- Supporting every country
- Collecting data without clear usage rights
- Hiding low-quality or incomplete regional data
- Making exact price promises
- Spending heavily on paid advertising

---

18. Revenue roadmap

Stage 1: Launch revenue

- Affiliate links
- Occasional direct sponsorship
- Voluntary supporters
- Newsletter sponsorship when audience permits

Stage 2: Growth revenue

- Display advertising
- Improved affiliate commerce
- Premium alerts
- Ad-free membership
- Sponsored reports
- Creator partnerships
- Chart licensing

Stage 3: Commercial revenue

- Professional subscriptions
- Historical data licensing
- API access
- Custom reports
- Retail dashboards
- White-labelled charts
- Enterprise data feeds
- Industry sponsorship

Revenue principles

- Forecasts must not be altered by commercial partners.
- Sponsored content must be clearly labelled.
- Retailer commission must not determine the price index.
- Ranking methodology must be published.
- Affiliate relationships must be disclosed.
- Data licensing must not remove public corrections.

---

19. Key performance indicators

Audience

- Monthly unique visitors
- Returning visitor rate
- Newsletter subscribers
- Newsletter open rate
- Direct traffic
- Search traffic
- Regional distribution

Product

- Regional selector usage
- Chart interaction
- Watchlist creation
- Alert creation
- Buy/wait interaction
- Product click-through rate
- Forecast archive usage

Data quality

- Product-match accuracy
- Successful daily imports
- Stale listings
- Retailer-source availability
- Index confidence
- Correction rate
- Data complaint rate

Editorial

- Forecast directional accuracy
- Forecast range accuracy
- Average forecast error
- Source diversity
- Correction response time
- Publication frequency

Commercial

- Affiliate click rate
- Affiliate conversion rate
- Revenue per visitor
- Revenue per newsletter subscriber
- Sponsor renewal rate
- Premium conversion
- B2B leads
- B2B recurring revenue

---

20. Risks and mitigations

Data licensing

Risk: A provider does not permit historical storage or commercial reuse.

Mitigation:

- Keep a data-source register.
- Record permitted uses.
- Prefer direct agreements and affiliate feeds.
- Do not build core indices on uncertain permissions.
- Obtain legal advice before commercial licensing.

Product matching errors

Risk: Similar products are incorrectly combined.

Mitigation:

- Prioritise manufacturer part numbers.
- Use confidence scores.
- Require human review for uncertain matches.
- Allow user error reports.
- Keep an audit trail.

Misleading regional comparisons

Risk: Tax and exchange-rate differences distort conclusions.

Mitigation:

- Show local prices and indexed movement separately.
- Clearly label tax treatment.
- Retain normalised comparison fields.
- Publish methodology.

Forecast credibility

Risk: Predictions appear overconfident or repeatedly fail.

Mitigation:

- Publish confidence.
- Show contradictory evidence.
- Use ranges.
- Record outcomes permanently.
- Revise the model based on performance.

Excessive scope

Risk: The team tries to support every market and category.

Mitigation:

- Maintain hard launch limits.
- Use geographic entry criteria.
- Require evidence before adding categories.
- Review scope at fixed intervals.

Dependence on search traffic

Risk: Algorithm changes reduce visitors.

Mitigation:

- Build the newsletter.
- Develop direct traffic.
- Create citation-worthy reports.
- Build professional products.
- Diversify revenue.

Commercial influence

Risk: Retailers or manufacturers pressure the team.

Mitigation:

- Separate editorial and commercial decisions.
- Publish methodology.
- Disclose sponsorship.
- Keep forecast history immutable.
- Reject payment for favourable analysis.

---

21. Governance and operating rhythm

Daily

- Monitor data imports.
- Review anomalies.
- Check major manufacturer and market events.
- Correct high-priority listing errors.

Weekly

- Publish global outlook.
- Update regional recommendations.
- Send newsletter.
- Review forecast assumptions.
- Review affiliate and audience performance.

Monthly

- Publish regional market summary.

- Review data-source health.
- Review forecast accuracy.
- Review product basket composition.
- Review roadmap priorities.
- Speak to users or partners.

Quarterly

- Publish forecast scorecard.
- Release an original market report.
- Review geographic expansion.
- Review commercial performance.
- Review methodology.
- Audit data licences and disclosures.

---

22. Immediate team actions

The team should begin with the following actions.

Product owner

- Finalise the product proposition.
- Confirm launch regions and categories.
- Schedule user interviews.
- Create the scope and decision log.

Technical lead

- Draft architecture.
- Select the data model.
- Define raw and derived data separation.
- Create the engineering backlog.

Data engineer

- Build the source register.
- Test retailer data.
- Create the first canonical product set.
- Prototype product matching.

Designer or front-end lead

- Produce homepage and regional-page wireframes.
- Define chart behaviour.

- Design confidence and quality indicators.
- Test mobile layouts.

Market editor

- Build the source hierarchy.
- Create event and forecast templates.
- Draft the initial global outlook.
- Establish corrections and review policies.

Commercial lead

- Identify affiliate programmes.
- Identify potential data partners.
- Interview system builders and retailers.
- Build the initial sponsor and B2B prospect list.

---

23. Decisions required before development starts

The team must make and document the following decisions:

1. Final working name
2. Initial four retail markets
3. Exact launch product baskets
4. Minimum number of retailers per region
5. Whether Germany represents the initial Eurozone index
6. How US sales tax will be displayed
7. Frequency of retail-price collection
8. Forecast time horizons
9. Confidence scale
10. Data-correction process
11. Affiliate-link policy
12. Open versus private code repository
13. Public availability of downloadable data
14. Minimum quality threshold for publishing an index
15. Definition of a successful private beta

---

24. Recommended launch definition

The MVP is ready for private beta when a user can:

1. Open the homepage and understand the global memory-market direction.
2. Select the United States, Eurozone, United Kingdom or Japan.

3. View a local-currency price history for DDR4, DDR5 or NVMe SSDs.
4. Compare regional movement using an indexed graph.
5. See major market events placed against price movement.
6. Understand why those events matter.
7. Read a buy/wait outlook with confidence and counterarguments.
8. Inspect the platform’s past forecasts.
9. Follow a disclosed affiliate link to an appropriate regional retailer.
10. Read exactly how the data and forecast were produced.

That is the first complete expression of the product.

Everything beyond it should be judged by whether it improves:

- Trust
- Data quality
- Repeat usage
- Consumer decisions
- Regional intelligence
- Revenue potential

The team should build the platform globally at the architectural and editorial level, while expanding retail coverage carefully enough that every supported region remains trustworthy.